

SF Salaries Exercise

January 18, 2018

1 SF Salaries Exercise

Welcome to a quick exercise for you to practice your pandas skills! We will be using the [SF Salaries Dataset](#) from Kaggle! Just follow along and complete the tasks outlined in bold below. The tasks will get harder and harder as you go along.

**** Import pandas as pd.****

```
In [6]: import pandas as pd
```

**** Read Salaries.csv as a dataframe called sal.****

```
In [7]: sal = pd.read_csv('Salaries.csv')
```

**** Check the head of the DataFrame. ****

```
In [8]: sal.head()
```

```
Out [8]:
```

	Id	EmployeeName	JobTitle	\
0	1	NATHANIEL FORD	GENERAL MANAGER-METROPOLITAN TRANSIT AUTHORITY	
1	2	GARY JIMENEZ	CAPTAIN III (POLICE DEPARTMENT)	
2	3	ALBERT PARDINI	CAPTAIN III (POLICE DEPARTMENT)	
3	4	CHRISTOPHER CHONG	WIRE ROPE CABLE MAINTENANCE MECHANIC	
4	5	PATRICK GARDNER	DEPUTY CHIEF OF DEPARTMENT, (FIRE DEPARTMENT)	

	BasePay	OvertimePay	OtherPay	Benefits	TotalPay	TotalPayBenefits	\
0	167411.18	0.00	400184.25	NaN	567595.43	567595.43	
1	155966.02	245131.88	137811.38	NaN	538909.28	538909.28	
2	212739.13	106088.18	16452.60	NaN	335279.91	335279.91	
3	77916.00	56120.71	198306.90	NaN	332343.61	332343.61	
4	134401.60	9737.00	182234.59	NaN	326373.19	326373.19	

	Year	Notes	Agency	Status
0	2011	NaN	San Francisco	NaN
1	2011	NaN	San Francisco	NaN
2	2011	NaN	San Francisco	NaN
3	2011	NaN	San Francisco	NaN
4	2011	NaN	San Francisco	NaN

**** Use the .info() method to find out how many entries there are.****

```
In [9]: sal.info() # 148654 Entries

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 148654 entries, 0 to 148653
Data columns (total 13 columns):
Id                148654 non-null int64
EmployeeName      148654 non-null object
JobTitle          148654 non-null object
BasePay           148045 non-null float64
OvertimePay       148650 non-null float64
OtherPay          148650 non-null float64
Benefits          112491 non-null float64
TotalPay          148654 non-null float64
TotalPayBenefits  148654 non-null float64
Year             148654 non-null int64
Notes             0 non-null float64
Agency           148654 non-null object
Status            0 non-null float64
dtypes: float64(8), int64(2), object(3)
memory usage: 14.7+ MB
```

What is the average BasePay ?

```
In [10]: sal['BasePay'].mean()
```

```
Out[10]: 66325.44884050643
```

**** What is the highest amount of OvertimePay in the dataset ? ****

```
In [11]: sal['OvertimePay'].max()
```

```
Out[11]: 245131.88
```

**** What is the job title of JOSEPH DRISCOLL ? Note: Use all caps, otherwise you may get an answer that doesn't match up (there is also a lowercase Joseph Driscoll). ****

```
In [12]: sal[sal['EmployeeName']=='JOSEPH DRISCOLL']['JobTitle']
```

```
Out[12]: 24    CAPTAIN, FIRE SUPPRESSION
         Name: JobTitle, dtype: object
```

**** How much does JOSEPH DRISCOLL make (including benefits)? ****

```
In [13]: sal[sal['EmployeeName']=='JOSEPH DRISCOLL']['TotalPayBenefits']
```

```
Out[13]: 24    270324.91
         Name: TotalPayBenefits, dtype: float64
```

**** What is the name of highest paid person (including benefits)?****

```
In [14]: sal[sal['TotalPayBenefits']== sal['TotalPayBenefits'].max()]
```

```
Out[14]:
```

	Id	EmployeeName	JobTitle	\
0	1	NATHANIEL FORD	GENERAL MANAGER-METROPOLITAN TRANSIT AUTHORITY	

	BasePay	OvertimePay	OtherPay	Benefits	TotalPay	TotalPayBenefits	\
0	167411.18	0.0	400184.25	NaN	567595.43	567595.43	

	Year	Notes	Agency	Status
0	2011	NaN	San Francisco	NaN

**** What is the name of lowest paid person (including benefits)? Do you notice something strange about how much he or she is paid? ****

```
In [15]: sal[sal['TotalPayBenefits']== sal['TotalPayBenefits'].min()]
```

```
Out[15]:
```

	Id	EmployeeName	JobTitle	BasePay	OvertimePay	\
148653	148654	Joe Lopez	Counselor, Log Cabin Ranch	0.0	0.0	

	OtherPay	Benefits	TotalPay	TotalPayBenefits	Year	Notes	\
148653	-618.13	0.0	-618.13	-618.13	2014	NaN	

	Agency	Status
148653	San Francisco	NaN

**** What was the average (mean) BasePay of all employees per year? (2011-2014) ? ****

```
In [16]: sal.groupby('Year').mean()['BasePay']
```

```
Out[16]: Year
2011    63595.956517
2012    65436.406857
2013    69630.030216
2014    66564.421924
Name: BasePay, dtype: float64
```

**** How many unique job titles are there? ****

```
In [17]: sal['JobTitle'].nunique()
```

```
Out[17]: 2159
```

**** What are the top 5 most common jobs? ****

```
In [18]: sal['JobTitle'].value_counts().head(5)
```

```
Out[18]: Transit Operator    7036
Special Nurse              4389
Registered Nurse           3736
Public Svc Aide-Public Works 2518
Police Officer 3           2421
Name: JobTitle, dtype: int64
```

**** How many Job Titles were represented by only one person in 2013? (e.g. Job Titles with only one occurrence in 2013?) ****

```
In [19]: sum(sal[sal['Year']==2013]['JobTitle'].value_counts() == 1)
```

```
Out[19]: 202
```

**** How many people have the word Chief in their job title? (This is pretty tricky) ****

```
In [20]: def chief_string(title):
          if 'chief' in title.lower():
              return True
          else:
              return False
```

```
In [21]: sum(sal['JobTitle'].apply(lambda x: chief_string(x)))
```

```
Out[21]: 477
```

**** Bonus: Is there a correlation between length of the Job Title string and Salary? ****

```
In [22]: sal['title_len'] = sal['JobTitle'].apply(len)
```

```
In [23]: sal[['title_len', 'TotalPayBenefits']].corr()
```

```
Out[23]:
```

	title_len	TotalPayBenefits
title_len	1.000000	-0.036878
TotalPayBenefits	-0.036878	1.000000

2 Great Job!